

Proj74ThesisReport_SuryaMunjal2

October 9, 2022

```
[1]: pip install lime
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-
wheels/public/simple/
Collecting lime
  Downloading lime-0.2.0.1.tar.gz (275 kB)
    |                               | 275 kB 28.4 MB/s
Requirement already satisfied: matplotlib in
/usr/local/lib/python3.7/dist-packages (from lime) (3.2.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages
(from lime) (1.21.6)
Requirement already satisfied: scipy in /usr/local/lib/python3.7/dist-packages
(from lime) (1.7.3)
Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages
(from lime) (4.64.1)
Requirement already satisfied: scikit-learn>=0.18 in
/usr/local/lib/python3.7/dist-packages (from lime) (1.0.2)
Requirement already satisfied: scikit-image>=0.12 in
/usr/local/lib/python3.7/dist-packages (from lime) (0.18.3)
Requirement already satisfied: tifffile>=2019.7.26 in
/usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime)
(2021.11.2)
Requirement already satisfied: PyWavelets>=1.1.1 in
/usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime) (1.3.0)
Requirement already satisfied: pillow!=7.1.0,!7.1.1,>=4.3.0 in
/usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime) (7.1.2)
Requirement already satisfied: imageio>=2.3.0 in /usr/local/lib/python3.7/dist-
packages (from scikit-image>=0.12->lime) (2.9.0)
Requirement already satisfied: networkx>=2.0 in /usr/local/lib/python3.7/dist-
packages (from scikit-image>=0.12->lime) (2.6.3)
Requirement already satisfied: pyparsing!=2.0.4,!2.1.2,!2.1.6,>=2.0.1 in
/usr/local/lib/python3.7/dist-packages (from matplotlib->lime) (3.0.9)
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.7/dist-
packages (from matplotlib->lime) (0.11.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.7/dist-packages (from matplotlib->lime) (1.4.4)
Requirement already satisfied: python-dateutil>=2.1 in
/usr/local/lib/python3.7/dist-packages (from matplotlib->lime) (2.8.2)
```

```

Requirement already satisfied: typing-extensions in
/usr/local/lib/python3.7/dist-packages (from
kiwisolver>=1.0.1->matplotlib->lime) (4.1.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-
packages (from python-dateutil>=2.1->matplotlib->lime) (1.15.0)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-
packages (from scikit-learn>=0.18->lime) (1.1.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.18->lime) (3.1.0)
Building wheels for collected packages: lime
  Building wheel for lime (setup.py) ... done
  Created wheel for lime: filename=lime-0.2.0.1-py3-none-any.whl size=283857
sha256=2d489edb3227bd90f833d68b3c71d84a3c171d0f85d9674deb5eb506bbc997ce
  Stored in directory: /root/.cache/pip/wheels/ca/cb/e5/ac701e12d365a08917bf4c61
71c0961bc880a8181359c66aa7
Successfully built lime
Installing collected packages: lime
Successfully installed lime-0.2.0.1

```

```

[2]: import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import sqlite3
from sqlalchemy import create_engine # database connection
import csv
import os
warnings.filterwarnings("ignore")
import datetime as dt
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
#from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.metrics import log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
#from sklearn.cross_validation import StratifiedKFold

```

```

from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier

from sklearn.model_selection import cross_val_score
from sklearn.linear_model import SGDClassifier
#from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import precision_recall_curve, auc, roc_curve
!pip install -q kaggle
import re
from bs4 import BeautifulSoup
from nltk.corpus import stopwords
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

from wordcloud import WordCloud, STOPWORDS
from os import path

import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import sqlite3
from sqlalchemy import create_engine # database connection
import csv
import os
warnings.filterwarnings("ignore")
import datetime as dt
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE

```

```

import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
#from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.metrics import log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
from sklearn.svm import SVC
#from sklearn.cross_validation import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier

from sklearn.model_selection import cross_val_score
from sklearn.linear_model import SGDClassifier
#from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import precision_recall_curve, auc, roc_curve
import numpy
from keras.datasets import imdb
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import LSTM
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence
# fix random seed for reproducibility
numpy.random.seed(7)

```

```

[3]: from google.colab import files
files.upload()

```

<IPython.core.display.HTML object>

Saving kaggle.json to kaggle (1).json

```
[3]: {'kaggle.json':  
      b'{"username":"suryamunjal","key":"376b73a1d8c76f470856f9ab7d5318e5"}'}
```

```
[4]: !mkdir ~/.kaggle
```

```
[5]: !cp kaggle.json ~/.kaggle/
```

```
[6]: !chmod 600 ~/.kaggle/kaggle.json
```

```
[7]: !kaggle datasets download -d akshayaki/fakenews
```

```
Downloading fakenews.zip to /content  
67% 24.0M/35.6M [00:00<00:00, 247MB/s]  
100% 35.6M/35.6M [00:00<00:00, 283MB/s]
```

```
[8]: !unzip fakenews.zip
```

```
Archive:  fakenews.zip  
  inflating: News.csv  
  
Reading the Data Set
```

```
[9]: df=pd.read_csv('News.csv')
```

```
[10]: df.head(2)
```

```
[10]:      Unnamed: 0      title \  
0      0  GRAPHIC RIOT VIDEOS EXPOSE THUGS ATTACKING ELD...  
1      1  BIG BROTHER: FEDS WANT YOUR DOCTOR TO WARN YOU...  
  
      text      subject \  
0  youngers these days are becoming so moist pic...  politics  
1  totally out of bounds! This is so wrong and so...  Government News  
  
      date Labels  
0  Sep 22, 2016  Fake  
1  Jun 26, 2015  Fake
```

```
[11]: df['combined']=df['title']+ " " + df["text"]
```

```
[12]: df['Labels']=df['Labels'].map({'Fake': 0, 'True': 1})
```

```
[13]: df=df[['combined','Labels']]
```

```
[14]: df.head(1)
```

```
[14]:      combined  Labels  
0  GRAPHIC RIOT VIDEOS EXPOSE THUGS ATTACKING ELD...      0
```

data cleaning

```
[15]: import nltk
from nltk.corpus import stopwords
nltk.download('stopwords')
stopwords = nltk.corpus.stopwords.words('english')

def decontracted(phrase):
    # specific
    phrase = re.sub(r"won't", "will not", phrase)
    phrase = re.sub(r"can't", "can not", phrase)

    # general
    phrase = re.sub(r"n't", " not", phrase)
    phrase = re.sub(r"\ 're", " are", phrase)
    phrase = re.sub(r"\ 's", " is", phrase)
    phrase = re.sub(r"\ 'd", " would", phrase)
    phrase = re.sub(r"\ 'll", " will", phrase)
    phrase = re.sub(r"\ 't", " not", phrase)
    phrase = re.sub(r"\ 've", " have", phrase)
    phrase = re.sub(r"\ 'm", " am", phrase)
    return phrase

def preproc(sentence):
    #from tqdm import tqdm
    # tqdm is for printing the status bar
    #for sentence in tqdm(df['combined'].values):
        sentence = re.sub(r"http\S+", "url", sentence) #removing urls with space
        sentence = BeautifulSoup(sentence, 'lxml').get_text() # removes tags like
        ↪ <br>
        sentence = decontracted(sentence)
        sentence = re.sub(r"\S*\d\S*", "", sentence).strip() # remove words with
        ↪ numbers
        sentence = re.sub(r'[^A-Za-z]+', ' ', sentence) ##remove spacial character:
        sentence = ' '.join(e.lower() for e in sentence.split() if e.lower() not in
        ↪ stopwords)

    return(sentence.strip())
```

[nltk_data] Downloading package stopwords to /root/nltk_data...

[nltk_data] Unzipping corpora/stopwords.zip.

```
[16]: df['Final_text']=df['combined'].apply(lambda x: preproc(str(x)))

#dropping the original combined column
df.drop(['combined'],inplace=True,axis=1)
```

```
[17]: df.head(2)
```

```
[17]:      Labels      Final_text
0      0  graphic riot videos expose thugs attacking eld...
1      0  big brother feds want doctor warn global warmi...
```

TRAIN AND TEST SPLIT

```
[18]: # split data into test and train

X_train, X_test, y_train, y_test = train_test_split(df['Final_text'], df.
↳Labels, test_size = 0.3)
```

```
[19]: X_train.shape
```

```
[19]: (31428,)
```

```
[20]: X_test.shape
```

```
[20]: (13470,)
```

preparing datasets for rnn models

```
[21]: import nltk
nltk.download('punkt')
from nltk.tokenize import word_tokenize, sent_tokenize

maxlen = -1
for i in df['Final_text']:
    tokens = nltk.word_tokenize(i)
    if(maxlen < len(tokens)):
        maxlen = len(tokens)
print("The maximum number of words in any document is =", maxlen)
```

[nltk_data] Downloading package punkt to /root/nltk_data...

[nltk_data] Unzipping tokenizers/punkt.zip.

The maximum number of words in any document is = 4963

```
[22]: list_of_words = []
for i in df['Final_text']:
    for j in i.split():
        list_of_words.append(j)
```

```
[23]: len(list_of_words)
```

```
[23]: 10833916
```

```
[24]: total=len(set(list_of_words))
      print("total no of unique words are",total)
```

total no of unique words are 105035

```
[25]: from tensorflow.keras.preprocessing.text import one_hot, Tokenizer
      voc_size=total #we could have giveb max like 5000
      #onehot_repr=[one_hot(words,total)for words in df['Final_text']]
      #onehot_repr[0]

      # Create a tokenizer to tokenize the words and create sequences of tokenized_
      ↪words
      tokenizer = Tokenizer(num_words = total)
      tokenizer.fit_on_texts(X_train)
      X_train = tokenizer.texts_to_sequences(X_train)

      X_test = tokenizer.texts_to_sequences(X_test)

      #X_train = tokenizer.sequences_to_matrix(train_sequences, mode='binary')
      #X_test = tokenizer.sequences_to_matrix(test_sequences, mode='binary')
```

```
[26]: #https://machinelearningknowledge.ai/
      ↪keras-tokenizer-tutorial-with-examples-for-fit_on_texts-texts_to_sequences-texts_to_matrix-
      ↪
      '''
      docs =['Machine Learning Knowledge',
              'Machine Learning and Deep Learning',
              'Deep Learning',
              'Artificial Intelligence']

      # create the tokenizer
      t = Tokenizer()

      t.fit_on_texts(docs)

      sequences = t.texts_to_sequences(docs)
      print(sequences)

      encoded_docs = t.sequences_to_matrix(sequences, mode='binary')
      print(encoded_docs)

      '''
```

```
[26]: "\ndocs =['Machine Learning Knowledge',\n              'Machine Learning and Deep Learning',\n              'Deep Learning',\n              'Artificial Intelligence']\n\n# create the tokenizer\nnt = Tokenizer()\n\nnt.fit_on_texts(docs)\n\nnsequences = t.texts_to_sequences(docs)\n\nnprint(sequences)\n\nnencoded_docs =
```


0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	3319	1598	11147	802	3138	1658	1	19982	19
4938	4464	14826	15	1	351	39	238	12241	3691	2895	5810
5753	4361	19	4438	19	2766	13799	5901	201	1119	19	2767
282	2755	927	1702	1755	394	1846	3555	43	1066	282	19983
1357	640	19	2766	557	10074	91	129	9702	5035	202	129
456	100	885	8211	800	1024	4464	223	410	1	1846	73
2698	132	282	84	37	264	1560	330	1030	4476	19001	2234
18163	6	3080	596	260	264	19	2766	251	1	524	4299
225	35	5810	5753	4361	1598	11147	802	3138	6800	899	8529
1	73	684	1978	800	1531	4	1	7730	4668	499	684
2755	3655	1388	909	4114	2813	400	802	3138	42880	539	388
1094	2939	388	1	2154	267	225	15	1	481	19	2767
12241	18	29	29	19	4438	1	771	141	4694	684	187
293	2847	1	4	117	4	19	4438	3033	3033	9918	18
37	2861	2704	516	29	151	107	76	1677	5431	438	370]

CREATING A LSTM MODEL

```
[29]: # create the model
embedding_vecor_length = 32
model = Sequential()
model.add(Embedding(total, embedding_vecor_length, input_length=max_length))
model.add(LSTM(8))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam',
    ↳metrics=['accuracy'])
print(model.summary())
#Refer: https://datascience.stackexchange.com/questions/10615/
    ↳number-of-parameters-in-an-lstm-model

#log loss is loss='binary_crossentropy' specifies that your model should
    ↳optimize the log loss for binary classification.
#metrics=['accuracy'] specifies that accuracy should be printed out, but log
    ↳loss is also printed out by default.
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 600, 32)	3361120
lstm (LSTM)	(None, 8)	1312

dense (Dense) (None, 1) 9

```
=====
Total params: 3,362,441
Trainable params: 3,362,441
Non-trainable params: 0
-----
None
```

```
[30]: history=model.fit(X_train, y_train, epochs=5, batch_size=64,validation_split=0.
      ↪3)
      # Final evaluation of the model
      scores = model.evaluate(X_test, y_test, verbose=0)
      print("Accuracy: %.2f%%" % (scores[1]*100))
```

```
Epoch 1/5
344/344 [=====] - 16s 25ms/step - loss: 0.2733 -
accuracy: 0.9179 - val_loss: 0.0993 - val_accuracy: 0.9785
Epoch 2/5
344/344 [=====] - 8s 23ms/step - loss: 0.0671 -
accuracy: 0.9859 - val_loss: 0.0689 - val_accuracy: 0.9833
Epoch 3/5
344/344 [=====] - 8s 23ms/step - loss: 0.0321 -
accuracy: 0.9937 - val_loss: 0.0647 - val_accuracy: 0.9837
Epoch 4/5
344/344 [=====] - 8s 23ms/step - loss: 0.0242 -
accuracy: 0.9945 - val_loss: 0.0576 - val_accuracy: 0.9846
Epoch 5/5
344/344 [=====] - 8s 23ms/step - loss: 0.0129 -
accuracy: 0.9978 - val_loss: 0.0606 - val_accuracy: 0.9836
Accuracy: 98.67%
```

```
[31]: print(history.history.keys())
```

```
dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])
```

```
[32]: print(history.history['val_loss'])
```

```
[0.09927283972501755, 0.0689343586564064, 0.06465856730937958,
0.057585202157497406, 0.060601186007261276]
```

```
[33]: print(history.history['val_accuracy'])
```

```
[0.9784706830978394, 0.9833492636680603, 0.9836674332618713, 0.9846218824386597,
0.9835613369941711]
```

```
[34]: print(scores)
```


```
[0.053384117782115936, 0.9867112040519714]
```

```
[35]: print(model.metrics_names)
```

```
['loss', 'accuracy']
```

1 so we are getting a test log loss of 0.05 in lstm

```
[36]: # plot train and validation loss
from matplotlib import pyplot
pyplot.plot(history.history['loss'])
pyplot.plot(history.history['val_loss'])
pyplot.title('model train vs validation loss')
pyplot.ylabel('loss')
pyplot.xlabel('epoch')
pyplot.legend(['train', 'validation'], loc='upper right')
pyplot.show()
```



```
[37]: from sklearn.metrics import confusion_matrix
y_pred = model.predict(X_test)
confusion_matrix = confusion_matrix(y_test, np rint(y_pred))
print(confusion_matrix)
```

```
from sklearn.metrics import classification_report
target_names = ['fake 0', 'real 1']
print(classification_report(y_test, np.rint(y_pred),target_names=target_names))
```

```
[[6995   99]
 [  80 6296]]

              precision    recall  f1-score   support

    fake 0           0.99        0.99        0.99        7094
    real 1           0.98        0.99        0.99        6376

 accuracy                   0.99        13470
 macro avg           0.99        0.99        0.99        13470
weighted avg           0.99        0.99        0.99        13470
```

```
[38]: print(confusion_matrix)
```

```
[[6995   99]
 [  80 6296]]
```

```
[39]: from sklearn.metrics import classification_report
target_names = ['fake 0', 'real 1']
print(classification_report(y_test, np.rint(y_pred),target_names=target_names))
```

```
              precision    recall  f1-score   support

    fake 0           0.99        0.99        0.99        7094
    real 1           0.98        0.99        0.99        6376

 accuracy                   0.99        13470
 macro avg           0.99        0.99        0.99        13470
weighted avg           0.99        0.99        0.99        13470
```

```
[40]: !pip install lime
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-
wheels/public/simple/
Requirement already satisfied: lime in /usr/local/lib/python3.7/dist-packages
(0.2.0.1)
Requirement already satisfied: scipy in /usr/local/lib/python3.7/dist-packages
(from lime) (1.7.3)
Requirement already satisfied: scikit-image>=0.12 in
/usr/local/lib/python3.7/dist-packages (from lime) (0.18.3)
Requirement already satisfied: scikit-learn>=0.18 in
/usr/local/lib/python3.7/dist-packages (from lime) (1.0.2)
Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages
```

(from lime) (4.64.1)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.7/dist-packages (from lime) (3.2.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from lime) (1.21.6)
Requirement already satisfied: networkx>=2.0 in /usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime) (2.6.3)
Requirement already satisfied: pillow!=7.1.0,!=7.1.1,>=4.3.0 in /usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime) (7.1.2)
Requirement already satisfied: PyWavelets>=1.1.1 in /usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime) (1.3.0)
Requirement already satisfied: tifffile>=2019.7.26 in /usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime) (2021.11.2)
Requirement already satisfied: imageio>=2.3.0 in /usr/local/lib/python3.7/dist-packages (from scikit-image>=0.12->lime) (2.9.0)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->lime) (2.8.2)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->lime) (1.4.4)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib->lime) (3.0.9)
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib->lime) (0.11.0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from kiwisolver>=1.0.1->matplotlib->lime) (4.1.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.1->matplotlib->lime) (1.15.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.18->lime) (3.1.0)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.18->lime) (1.1.0)

```
[41]: from lime.lime_text import LimeTextExplainer
class_names=['0','1']
explainer= LimeTextExplainer(class_names=class_names)
def predict_proba(arr):
    processed=[]
    for i in arr:
        processed.append(preproc(i))
    list_tokenized_ex = tokenizer.texts_to_sequences(processed)
    Ex = pad_sequences(list_tokenized_ex, maxlen=max_length)
    pred=model.predict(Ex)
    returnable=[]
    for i in pred:
        temp=i[0]
```

```

    returnable.append(np.array([1-temp,temp])) #I would recommend rounding temp
    ↪ and 1-temp off to 2 places
    return np.array(returnable)

```

```

[42]: data=pd.read_csv('News.csv')
      data['combined']=data['title']+" "+data['text']
      data.head(1)

```

```

[42]: Unnamed: 0                                     title \
0          0  GRAPHIC RIOT VIDEOS EXPOSE THUGS ATTACKING ELD...

                                     text  subject          date \
0  youngers these days are becoming so moist pic...  politics  Sep 22, 2016

Labels                                     combined
0  Fake  GRAPHIC RIOT VIDEOS EXPOSE THUGS ATTACKING ELD...

```

```

[43]: print("Actual",data['Labels'][666])
      explainer.explain_instance(data['combined'][666],predict_proba).
      ↪ show_in_notebook(text=True)

```

Actual Fake

<IPython.core.display.HTML object>

```

[44]: print("Actual",data['Labels'][4437])
      explainer.explain_instance(data['combined'][4437],predict_proba).
      ↪ show_in_notebook(text=True)

```

Actual True

<IPython.core.display.HTML object>

2 LETS CREATE A BIDIRECTIONAL RNN MODEL

```

[45]: ## Creating model
      from tensorflow.keras.layers import Bidirectional
      from tensorflow.keras.layers import Dropout
      # create the model

      embedding_vecor_length = 32
      model1 = Sequential()
      model1.add(Embedding(total, embedding_vecor_length, input_length=max_length))
      model1.add(Bidirectional(LSTM(8)))
      model1.add(Dropout(0.3))

```

```

model1.add(Dense(1, activation='sigmoid'))
model1.compile(loss='binary_crossentropy', optimizer='adam',
↳metrics=['accuracy'])
print(model1.summary())

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 600, 32)	3361120
bidirectional (Bidirectional) 1)	(None, 16)	2624
dropout (Dropout)	(None, 16)	0
dense_1 (Dense)	(None, 1)	17

Total params: 3,363,761
 Trainable params: 3,363,761
 Non-trainable params: 0

None

```

[46]: history=model1.fit(X_train, y_train, epochs=5, batch_size=64, validation_split=0.
↳3)
# Final evaluation of the model
scores = model1.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1]*100))

```

Epoch 1/5

344/344 [=====] - 27s 69ms/step - loss: 0.2967 - accuracy: 0.9174 - val_loss: 0.1252 - val_accuracy: 0.9722

Epoch 2/5

344/344 [=====] - 16s 48ms/step - loss: 0.0874 - accuracy: 0.9850 - val_loss: 0.0816 - val_accuracy: 0.9773

Epoch 3/5

344/344 [=====] - 14s 41ms/step - loss: 0.0576 - accuracy: 0.9886 - val_loss: 0.0649 - val_accuracy: 0.9822

Epoch 4/5


344/344 [=====] - 14s 41ms/step - loss: 0.0393 - accuracy: 0.9933 - val_loss: 0.0878 - val_accuracy: 0.9736

Epoch 5/5

344/344 [=====] - 14s 41ms/step - loss: 0.0464 - accuracy: 0.9903 - val_loss: 0.0848 - val_accuracy: 0.9750

Accuracy: 97.49%


```
[47]: # plot train and validation loss
from matplotlib import pyplot
pyplot.plot(history.history['loss'])
pyplot.plot(history.history['val_loss'])
pyplot.title('model train vs validation loss')
pyplot.ylabel('loss')
pyplot.xlabel('epoch')
pyplot.legend(['train', 'validation'], loc='upper right')
pyplot.show()
```



```
[48]: from sklearn.metrics import confusion_matrix
y_pred1 = model1.predict(X_test)
confusion_matrix1 = confusion_matrix(y_test, np.rint(y_pred1))
print(confusion_matrix1)
from sklearn.metrics import classification_report
target_names = ['fake 0', 'real 1']
print(classification_report(y_test, np.rint(y_pred1), target_names=target_names))
```

```
[[6933 161]
 [177 6199]]
```

	precision	recall	f1-score	support
fake 0	0.98	0.98	0.98	7094
real 1	0.97	0.97	0.97	6376

accuracy			0.97	13470
macro avg	0.97	0.97	0.97	13470
weighted avg	0.97	0.97	0.97	13470

```
[49]: from lime.lime_text import LimeTextExplainer
class_names=['0','1']
explainer= LimeTextExplainer(class_names=class_names)
def predict_proba(arr):
    processed=[]
    for i in arr:
        processed.append(preproc(i))
    list_tokenized_ex = tokenizer.texts_to_sequences(processed)
    Ex = pad_sequences(list_tokenized_ex, maxlen=max_length)
    pred=model.predict(Ex)
    returnable=[]
    for i in pred:
        temp=i[0]
        returnable.append(np.array([1-temp,temp])) #I would recommend rounding temp
    →and 1-temp off to 2 places
    return np.array(returnable)
```

```
[50]: from lime.lime_text import LimeTextExplainer
class_names=['0','1']
explainer1= LimeTextExplainer(class_names=class_names)
def predict_proba1(arr):
    processed=[]
    for i in arr:
        processed.append(preproc(i))
    list_tokenized_ex = tokenizer.texts_to_sequences(processed)
    Ex = pad_sequences(list_tokenized_ex, maxlen=max_length)
    pred=model1.predict(Ex)
    returnable=[]
    for i in pred:
        temp=i[0]
        returnable.append(np.array([1-temp,temp])) #I would recommend rounding temp
    →and 1-temp off to 2 places
    return np.array(returnable)
```

```
[51]: print("Actual",data['Labels'][769])
explainer1.explain_instance(data['combined'][669],predict_proba1).
→show_in_notebook(text=True)
```

Actual Fake

<IPython.core.display.HTML object>

```
[53]: import joblib
      joblib.dump(model, 'model.pkl')
      joblib.dump(tokenizer, 'tokenizer.pkl')
```

WARNING:absl:Found untraced functions such as lstm_cell_layer_call_fn, lstm_cell_layer_call_and_return_conditional_losses while saving (showing 2 of 2). These functions will not be directly callable after loading.

WARNING:absl:<keras.layers.recurrent.LSTMCell object at 0x7f52355e3e10> has the same name 'LSTMCell' as a built-in Keras object. Consider renaming <class 'keras.layers.recurrent.LSTMCell'> to avoid naming conflicts when loading with `tf.keras.models.load_model`. If renaming is not possible, pass the object in the `custom_objects` parameter of the load function.

```
[53]: ['tokenizer.pkl']
```

```
[ ]:
```